

Evaluation ideal and variety for a trio of error independent binary classifiers

Introduction

This notebook will detail the algebraic geometry computations that take us from the “evaluation ideal” created from the voting patterns of a trio of binary classifiers to the “evaluation variety”. An evaluation ideal is a set of polynomials connecting observable voting pattern frequencies by the classifiers to unknown sample statistics of the ground truth that are our evaluation goal. We want to “grade” the classifiers using only the frequencies of their voting patterns.

That “grade” exists in sample statistics space. The test has already been taken. We have the decisions of the judges. We are faced with the task of grading them now. Not in the future, not in the past. This is another example of how the task of evaluation is much simpler than that of training. We have to estimate something that already exists, if you will. And there is only one time we have to do it. Training is much harder. You must create judges that, in the future, will behave correctly. And they have to do it many times. The task of evaluation is trivial in comparison. Why have we not conquered this much simpler space of the whole enterprise of learning?

Algebraic geometry of three error independent binary classifiers

The mathematics of algebraic evaluation is algebraic geometry. Every algebraic evaluation problem can be stated as a polynomial system relating observable decision events to unknown sample statistics. Here we are going to define that polynomial system assuming that the classifiers made errors independently on the sample. This is “the spherical cow” of Evaluation Land - the simplifying assumption that allows you to proceed forward and carry out computations that give you insight into the original problem. Workers in Training Land also have a preferred spherical cow - “consider an identically, independently drawn sample”. It may take some getting used to this new cow if you are a new visitor from Training Land.

The evaluation ideal of three error independent binary classifiers

Algebraic geometry is mainly the study of the connection between sets of polynomials and geometric objects in the variable space of those polynomials. The sets of polynomials are called “polynomial

ideals". A set of linear equations is also a polynomial ideal. We define the "evaluation ideal" of our evaluation to be,

In[3165]:=

```
Clear[MakeIndependentVotingIdeal]
MakeIndependentVotingIdeal[{i_, j_, k_}] :=
{Pα Pi,α Pj,α Pk,α + (1 - Pα) (1 - Pi,β) (1 - Pj,β) (1 - Pk,β) - fα,α,α,
Pα Pi,α Pj,α (1 - Pk,α) + (1 - Pα) (1 - Pi,β) (1 - Pj,β) Pk,β - fα,α,β,
Pα Pi,α (1 - Pj,α) Pk,α + (1 - Pα) (1 - Pi,β) Pj,β (1 - Pk,β) - fα,β,α,
Pα Pi,α (1 - Pj,α) (1 - Pk,α) + (1 - Pα) (1 - Pi,β) Pj,β Pk,β - fα,β,β,
Pα (1 - Pi,α) Pj,α Pk,α + (1 - Pα) Pi,β (1 - Pj,β) (1 - Pk,β) - fβ,α,α,
Pα (1 - Pi,α) Pj,α (1 - Pk,α) + (1 - Pα) Pi,β (1 - Pj,β) Pk,β - fβ,α,β,
Pα (1 - Pi,α) (1 - Pj,α) Pk,α + (1 - Pα) Pi,β Pj,β (1 - Pk,β) - fβ,β,α,
Pα (1 - Pi,α) (1 - Pj,α) (1 - Pk,α) + (1 - Pα) Pi,β Pj,β Pk,β - fβ,β,β}
```

One convention in algebraic geometry may bother you. Ultimately we are interested in the geometrical object these polynomials define in the finite space needed for evaluating three independent binary classifiers. We want to consider the points in sample statistics space where all these equations are zero. We are really interested in these equations,

In[3167]:=

```
MakeIndependentVotingIdeal[{1, 2, 3}] // Map[(# == 0) &, #] &
```

Out[3167]=

```
{Pα P1,α P2,α P3,α + (1 - Pα) (1 - P1,β) (1 - P2,β) (1 - P3,β) - fα,α,α == 0,
Pα P1,α P2,α (1 - P3,α) + (1 - Pα) (1 - P1,β) (1 - P2,β) P3,β - fα,α,β == 0,
Pα P1,α (1 - P2,α) P3,α + (1 - Pα) (1 - P1,β) P2,β (1 - P3,β) - fα,β,α == 0,
Pα P1,α (1 - P2,α) (1 - P3,α) + (1 - Pα) (1 - P1,β) P2,β P3,β - fα,β,β == 0,
Pα (1 - P1,α) P2,α P3,α + (1 - Pα) P1,β (1 - P2,β) (1 - P3,β) - fβ,α,α == 0,
Pα (1 - P1,α) P2,α (1 - P3,α) + (1 - Pα) P1,β (1 - P2,β) P3,β - fβ,α,β == 0,
Pα (1 - P1,α) (1 - P2,α) P3,α + (1 - Pα) P1,β P2,β (1 - P3,β) - fβ,β,α == 0,
Pα (1 - P1,α) (1 - P2,α) (1 - P3,α) + (1 - Pα) P1,β P2,β P3,β - fβ,β,β == 0}
```

But it is a pain to carry all these equal to zero notation around. So we drop it, and prefer to work with,

In[3168]:=

```
MakeIndependentVotingIdeal[{1, 2, 3}]
```

Out[3168]=

```
{Pα P1,α P2,α P3,α + (1 - Pα) (1 - P1,β) (1 - P2,β) (1 - P3,β) - fα,α,α,
Pα P1,α P2,α (1 - P3,α) + (1 - Pα) (1 - P1,β) (1 - P2,β) P3,β - fα,α,β,
Pα P1,α (1 - P2,α) P3,α + (1 - Pα) (1 - P1,β) P2,β (1 - P3,β) - fα,β,α,
Pα P1,α (1 - P2,α) (1 - P3,α) + (1 - Pα) (1 - P1,β) P2,β P3,β - fα,β,β,
Pα (1 - P1,α) P2,α P3,α + (1 - Pα) P1,β (1 - P2,β) (1 - P3,β) - fβ,α,α,
Pα (1 - P1,α) P2,α (1 - P3,α) + (1 - Pα) P1,β (1 - P2,β) P3,β - fβ,α,β,
Pα (1 - P1,α) (1 - P2,α) P3,α + (1 - Pα) P1,β P2,β (1 - P3,β) - fβ,β,α,
Pα (1 - P1,α) (1 - P2,α) (1 - P3,α) + (1 - Pα) P1,β P2,β P3,β - fβ,β,β}
```

Dropping the notation has no effect. Algebraic manipulations of the above set (multiplying them together, etc.) would be equivalent to polynomials of zero for points that satisfied the input set.

One “forest for the trees” note: the above ideal is one many possible ones. Algebraic evaluation is very much like data streaming algorithms. You are creating a sketch of the decisions by an ensemble of noisy judges. In the case of binary classification considered here, that “data sketch” is the frequency of their item-by-item voting patterns. Other polynomial systems are possible even for independent binary classifiers. We could be trying to evaluate sample statistics that look at how judges evaluated two different sample items, for example.

The evaluation variety of three error independent binary classifiers

The goal of algebraic evaluation is to obtain “grades” for the noisy judges on unlabeled data. We know the types of grades we want. We want their exact grades. Not spreads of where we think their grade is - no probabilistic solutions! That is not quite what we get in algebraic evaluation. The mathematical object you get as a grade is actually a geometric object in sample statistics space. The true grade lies on that geometric object. Mathematicians call the geometrical objects defined by polynomials - varieties. You will see that for error-independent binary classifiers that geometric object is almost what we want - a collection of two points. For correlated classifiers those points “bloom out”. It is an unsolved problem in algebraic evaluation to characterize the surface for correlated classifiers - but it does exist! The evaluation ideal for any set of correlated classifiers can be written down and it defines, by construction, an evaluation variety that is guaranteed to contain the true evaluation point for the classifiers. All that is left when you build the evaluation ideal is to figure out what that surface is and whether it would be useful to your AI safety task.

A quick look at evaluation

Before continuing to present the formalism of algebraic evaluation, let’s do a quick end run to our goal - the grade for noisy binary classifiers. I’ll use the UCI Adult run that can be found in the Python code - AlgebraicEvaluation.py. The input to algebraic evaluation in our current application is the frequency of voting patterns by the classifiers. Here is how it looked for a single run of three binary classifiers on the UCI Adult dataset.

In[3169]:=

```
singleEvaluationUCIAdult =
<|0 → <|{0, 0, 0} → 715, {0, 0, 1} → 161, {0, 1, 0} → 2406, {0, 1, 1} → 455,
  {1, 0, 0} → 290, {1, 0, 1} → 94, {1, 1, 0} → 1335, {1, 1, 1} → 231|>,
  1 → <|{0, 0, 0} → 271, {0, 0, 1} → 469, {0, 1, 0} → 3395, {0, 1, 1} → 7517,
  {1, 0, 0} → 272, {1, 0, 1} → 399, {1, 1, 0} → 6377, {1, 1, 1} → 12 455|>|>
```

Out[3169]=

```
<|0 → <|{0, 0, 0} → 715, {0, 0, 1} → 161, {0, 1, 0} → 2406, {0, 1, 1} → 455,
  {1, 0, 0} → 290, {1, 0, 1} → 94, {1, 1, 0} → 1335, {1, 1, 1} → 231|>,
  1 → <|{0, 0, 0} → 271, {0, 0, 1} → 469, {0, 1, 0} → 3395, {0, 1, 1} → 7517,
  {1, 0, 0} → 272, {1, 0, 1} → 399, {1, 1, 0} → 6377, {1, 1, 1} → 12 455|>|>
```

This is not a randomly selected run of binary classifiers on the UCI Adult dataset. It is being used for various reasons. It was engineered to be as close to error independence as possible. This will be discussed more later. For now, let's verify that, in fact, the classifiers are near error independence in this sample.

In[3220]:=

```
evaluationGroundTruth = GTClassifiers[singleEvaluationUCIAdult]
```

Out[3220]=

$$\left\langle \left| P_{\alpha} \rightarrow \frac{5687}{36\,842}, P_{1,\alpha} \rightarrow \frac{3737}{5687}, P_{2,\alpha} \rightarrow \frac{1260}{5687}, P_{3,\alpha} \rightarrow \frac{4746}{5687}, P_{1,\beta} \rightarrow \frac{6501}{10\,385}, \right. \right.$$

$$P_{2,\beta} \rightarrow \frac{29\,744}{31\,155}, P_{3,\beta} \rightarrow \frac{4168}{6231}, \Gamma_{1,2,\alpha} \rightarrow \frac{273\,192}{32\,341\,969}, \Gamma_{1,3,\alpha} \rightarrow \frac{13\,325}{32\,341\,969},$$

$$\Gamma_{2,3,\alpha} \rightarrow -\frac{264\,525}{32\,341\,969}, \Gamma_{1,2,\beta} \rightarrow \frac{2\,204\,576}{323\,544\,675}, \Gamma_{1,3,\beta} \rightarrow -\frac{79\,682}{12\,941\,787},$$

$$\left. \Gamma_{2,3,\beta} \rightarrow \frac{94\,508}{38\,825\,361}, \Gamma_{1,2,3,\alpha} \rightarrow \frac{452\,568\,508}{183\,928\,777\,703}, \Gamma_{1,2,3,\beta} \rightarrow -\frac{27\,265\,589}{134\,400\,457\,995} \right| \rangle$$

The evaluation ground truth is what we want. It was calculated above by cheating - we have the by-label counts so we can easily compute ALL the sample statistics required to explain exactly the frequency patterns we observe when we DO NOT have the knowledge of the true labels. Note that the evaluation ground truth are integer ratios. The exact sample statistics for evaluation are a subset of the real numbers. This is crucial. This allows algebraic evaluators to get closer to the true grades. Integer ratios are also in the field of algebraic numbers. But the field of algebraic numbers is less dense than reals!

This evaluation also reminds the reader of why evaluation is easier than training. There are no unknown unknowns. Evaluation computes sample statistics. The space of sample statistics required to explain observable voting patterns is finite and complete. You may not know what all these sample statistics are, but they are all you would need to know to describe the frequency of their observed decisions.

It is hard to compare integer ratios so let's get the floating point approximation to the evaluation to confirm the claim that these classifiers are nearly error independent.

In[3221]:=

evaluationGroundTruth // N

Out[3221]=

$$\langle | P_{\alpha} \rightarrow 0.154362, P_{1,\alpha} \rightarrow 0.657113, P_{2,\alpha} \rightarrow 0.221558, P_{3,\alpha} \rightarrow 0.834535, \\ P_{1,\beta} \rightarrow 0.625999, P_{2,\beta} \rightarrow 0.95471, P_{3,\beta} \rightarrow 0.668913, \Gamma_{1,2,\alpha} \rightarrow 0.00844698, \\ \Gamma_{1,3,\alpha} \rightarrow 0.000412003, \Gamma_{2,3,\alpha} \rightarrow -0.008179, \Gamma_{1,2,\beta} \rightarrow 0.00681382, \Gamma_{1,3,\beta} \rightarrow -0.00615695, \\ \Gamma_{2,3,\beta} \rightarrow 0.00243418, \Gamma_{1,2,3,\alpha} \rightarrow 0.00246056, \Gamma_{1,2,3,\beta} \rightarrow -0.000202868 | \rangle$$

One can see that all the pair error correlation terms ($\Gamma_{i,j,\text{label}}$) are less that 1% absolute.